

CS553 Programming Assignment #1

Benchmarking

Instructions:

- *Assigned date: Monday February 19th, 2018*
- *Due date: 11:59PM on Friday March 9th, 2018*
- *Maximum Points: 100%*
- *This programming assignment must be done individually*
- *Please post your questions to the Piazza forum*
- *Only a softcopy submission is required; it must be submitted through GIT*
- *Late submission will be penalized at 20% per day (beyond the 4-day late pass).*

1 Your Assignment

This project aims to teach you how to benchmark different parts of a computer system, from the CPU, memory, disk, and network. You can be creative with this project. You are free to use any of the following programming languages (C, C++, Java, Python) and abstractions (PThreads, Sockets) that might be needed. Other programming languages will not be allowed due to the increased complexity in grading; do not write code that relies on complex libraries (e.g. boost), as those will simplify certain parts of your assignments, and will increase complexity in grading.

You can use any Linux system for your development, but you must use the virtual cluster found at 129.114.33.105 for the formal testing and evaluation. This virtual cluster is running on the Chameleon testbed [<https://www.chameleoncloud.org>]; more information about the hardware in this testbed can be found at <https://www.chameleoncloud.org/about/hardware-description/>, under Standard Cloud Units. Even more details can be found at <https://www.chameleoncloud.org/user/discovery/>, choose “Compute”, then Click the “View” button.

You have been created accounts on this virtual cluster for this assignment; if you have not received your accounts information, reach out to the TAs for this information. The computer system found at 129.114.33.105 is the cluster login node. You can connect to this login node with ssh. Do not use the login nodes for anything other than ssh from your system to the cluster, and other light commands (e.g. file system operations, compilation, code editing, etc). Once on the login node, you must use the Slurm job management system to submit jobs to run your benchmarks. You must use GIT for source control throughout your assignment; we will use GIT to also collect assignment submissions.

The cluster has a NFS setup and configured at /exports/home/userid with a 20GB quota per user. This directory can be used to store data that you want to access across the virtual cluster. There are two additional storage resources that are considered node local storage: SSD (/tmp with 20GB quota, it is cleaned up after every job) and RAMDisk (/dev/shm with variable limited space, it is cleaned up after every job). Please pay attention to where you must run your benchmarks from (e.g. NFS, SSD, or RAMDisk).

In this project, you need to design a benchmarking program that covers the four system components: processor, memory, disk, and network. You will perform strong scaling studies, unless otherwise noted; this means you will set the amount of work (e.g. the number of instructions or the amount of data to evaluate in your benchmark), and reduce the amount of work per thread as you increase the number of threads. The TAs will compile (with the help of make, ant, or maven) and test your code on this virtual

cluster. If your code does not compile and the TAs cannot run your project, you will get 0 for the assignment.

1. Processor:

- a. Implement: MyCPUBench benchmark
- b. Workload: 1 trillion arithmetic (quarter precision, half precision, single precision, double precision) operations
 - i. QP: quarter precision operations compute on 1-byte char data types
 - ii. HP: half precision operations compute on 2-byte short data types
 - iii. SP: single precision operations compute on 4-byte integers data types
 - iv. DP: double precision operations compute on 8-byte double data types
- c. Concurrency: 1 thread, 2 threads, 4 threads
- d. Measure: processor speed, in terms of operations per second; report data in GigaOPS, giga operations (10^9) per second
- e. Run the HPL benchmark from the Linpack suite (<http://en.wikipedia.org/wiki/LINPACK>) and report the best performance achieved using double precision floating point; make sure to run Linpack across all cores and to use a problem size that is large enough to consume $\frac{3}{4}$ of the available memory of your testing node (for this benchmark, you do not have to compute 1 trillion arithmetic computations, but the problem size you set to fill $\frac{3}{4}$ of the available memory will dictate the workload size). You can download the Linpack benchmark suite from <https://software.intel.com/en-us/articles/intel-mkl-benchmarks-suite> (binaries of the HPL benchmarks are available at this link, there is no need to compile the benchmark); run the HPL without MPI benchmark from the suite. Make sure to tune the HPL benchmark in HPL.dat (see <http://www.netlib.org/benchmark/hpl/tuning.html> for more information on how to interpret the HPL.dat configuration)
- f. Compute the theoretical peak performance of your processor. What efficiency do you achieve compared to the theoretical performance? Compare and contrast your performance to that achieved by your benchmark, HPL, and the theoretical peak performance.
- g. Fill in the table 1 below for Processor Performance:

Workload	Concurrency	MyCPUBench Measured Ops/Sec (GigaOPS)	HPL Measured Ops/Sec (GigaOPS)	Theoretical Ops/Sec (GigaOPS)	MyCPUBench Efficiency (%)	HPL Efficiency (%)
QP	1		N/A			N/A
QP	2		N/A			N/A
QP	4		N/A			N/A
HP	1		N/A			N/A
HP	2		N/A			N/A
HP	4		N/A			N/A
SP	1		N/A			N/A
SP	2		N/A			N/A
SP	4		N/A			N/A
DP	1					
DP	2					
DP	4					

2. Memory:

- Implement: MyRAMBench benchmark; *hint: you are unlikely going to be able to do this benchmark in Java or other high level languages, while C/C++ is a natural language to implement this benchmark; make sure to measure the performance of your memory and not your processor caches*
- Workload: 1GB data; operate over it 100X times with various access patterns (RWS, RWR) and various block sizes (1KB, 1MB, 10MB)
 - RWS: read+write (e.g. memcpy) with sequential access pattern
 - RWR: read+write (e.g. memcpy) with random access pattern
- Concurrency: 1 thread, 2 threads, 4 threads
- Measure:
 - throughput, in terms of bytes per second; report data in GB/sec, gigabytes (10^9) per second; these experiments should be conducted over 100GB of data
 - latency (read+write 1 byte of data), in terms of time per access; report data in us, microseconds; limit experiment to 100 million operations
- Run the Memory benchmark pmbw (<https://panthema.net/2013/pmbw/>) with 1, 2, and 4 threads, and measure the memory sub-system performance
- Compute the theoretical bandwidth and latency of your memory. What efficiency do you achieve compared to the theoretical performance? Compare and contrast your performance to that achieved by pmbw, and to the theoretical performance.
- Fill in the table 2 below for Memory Throughput:

Work-load	Con-currency	Block Size	MyRAMBench Measured Throughput (GB/sec)	pmbw Measured Throughput (GB/sec)	Theoretical Throughput (GB/sec)	MyRAMBench Efficiency (%)	pmbw Efficiency (%)
RWS	1	1KB					
RWS	1	1MB					
RWS	1	10MB					
RWS	2	1KB					
RWS	2	1MB					
RWS	2	10MB					
RWS	4	1KB					
RWS	4	1MB					
RWS	4	10MB					
RWR	1	1KB					
RWR	1	1MB					
RWR	1	10MB					
RWR	2	1KB					
RWR	2	1MB					
RWR	2	10MB					
RWR	4	1KB					
RWR	4	1MB					
RWR	4	10MB					

- Fill in the table 3 below for Memory Latency:

Work-load	Con-currency	Block Size	MyRAMBench Measured Latency (us)	pmbw Measured Latency (us)	Theoretical Latency (us)	MyRAMBench Efficiency (%)	pmbw Efficiency (%)
RWS	1	1B					
RWS	2	1B					
RWS	4	1B					
RWR	1	1B					
RWR	2	1B					
RWR	4	1B					

3. Disk:

- Implement: MyDiskBench benchmark; **Hint:** there are multiple ways to read and write to disk, explore the different APIs, and pick the fastest one out of all them; also make sure you are measuring the speed of your disk and not your memory (you may need to flush your disk cache managed by the OS)
- Workload: 10GB data; operate over it 1X times with various access patterns (RWS, RWR) and various block sizes (1MB, 10MB, 100MB)
 - RS: read with sequential access pattern
 - WS: write with sequential access pattern
 - RR: read with random access pattern
 - WS: write with random access pattern
- Concurrency: 1 thread, 2 threads, 4 threads
- Measure:
 - throughput, in terms of bytes per second; report data in MB/sec, megabytes (10^6) per second; these experiments should be conducted over 10GB of data
 - latency (read or write 1KB of data), in terms of time per access; report data in ms, milliseconds and in IOPS (I/O operations per second); limit experiment to 1 million operations (1GB data)
- Run the Disk benchmark IOZone benchmark (<http://www.iozone.org/>) with 1, 2, and 4 threads, and measure the disk performance
- Compute the theoretical bandwidth and latency of your disk. What efficiency do you achieve compared to the theoretical performance? Compare and contrast your performance to that achieved by IOZone, and to the theoretical performance.
- Fill in the table 4 below for Disk Throughput:

Work-load	Con-currency	Block Size	MyDiskBench Measured Throughput (MB/sec)	IOZone Measured Throughput (MB/sec)	Theoretical Throughput (MB/sec)	MyDiskBench Efficiency (%)	IOZone Efficiency (%)
RS	1	1MB					
RS	1	10MB					
RS	1	100MB					
RS	2	1MB					
RS	2	10MB					
RS	2	100MB					
RS	4	1MB					

RS	4	10MB					
RS	4	100MB					
WS	1	1MB					
WS	1	10MB					
WS	1	100MB					
WS	2	1MB					
WS	2	10MB					
WS	2	100MB					
WS	4	1MB					
WS	4	10MB					
WS	4	100MB					
RR	1	1MB					
RR	1	10MB					
RR	1	100MB					
RR	2	1MB					
RR	2	10MB					
RR	2	100MB					
RR	4	1MB					
RR	4	10MB					
RR	4	100MB					
WR	1	1MB					
WR	1	10MB					
WR	1	100MB					
WR	2	1MB					
WR	2	10MB					
WR	2	100MB					
WR	4	1MB					
WR	4	10MB					
WR	4	100MB					

h. Fill in the table 5 below for Disk Latency (measured in ms):

Work-load	Con-currency	Block Size	MyDiskBench Measured Latency (ms)	IOZone Measured Latency (ms)	Theoretical Latency (ms)	MyDiskBench Efficiency (%)	IOZone Efficiency (%)
RR	1	1KB					
RR	2	1KB					
RR	4	1KB					
RR	8	1KB					
RR	16	1KB					
RR	32	1KB					
RR	64	1KB					
RR	128	1KB					
WR	1	1KB					
WR	2	1KB					
WR	4	1KB					

WR	8	1KB					
WR	16	1KB					
WR	32	1KB					
WR	64	1KB					
WR	128	1KB					

- i. Fill in the table 6 below for Disk Latency (measured in IOPS); this data can be collected at the same time as the data collected from table 5:

Work-load	Con-currency	Block Size	MyDiskBench Measured IOPS	IOZone Measured IOPS	Theoretical IOPS	MyDiskBench Efficiency (%)	IOZone Efficiency (%)
RR	1	1KB					
RR	2	1KB					
RR	4	1KB					
RR	8	1KB					
RR	16	1KB					
RR	32	1KB					
RR	64	1KB					
RR	128	1KB					
WR	1	1KB					
WR	2	1KB					
WR	4	1KB					
WR	8	1KB					
WR	16	1KB					
WR	32	1KB					
WR	64	1KB					
WR	128	1KB					

4. Network:

- Implement: MyNETBench benchmark (must support both client and server functionality in the same program; multi-threaded support must exist at both client and server; **hint:** you are going to need two nodes for these experiments, one for the server (receiver) and one for the client (sender))
- Workload: 1GB data; operate over it 100X times with various block sizes (1KB, 32KB)
- Protocols: evaluate both the TCP and UDP protocol
- Concurrency: 1 thread, 2 threads, 4 threads, 8 threads
- Measure:
 - throughput, in terms of bytes per second; report data in Mb/sec, megabits (10^6) per second; these experiments should be conducted over 100GB of data to send from a client and be received at a server; the data should be kept in memory in a 1GB piece of memory on the client, and it should be stored in a 1GB piece of memory on the server that it overwrites as more data is received
 - latency (ping-pong 1 byte of data), in terms of time per RTT (round-trip-time); report data in ms, milliseconds; limit experiment to 1 million operations

- f. Run the Network benchmark iperf (<http://en.wikipedia.org/wiki/Iperf>) with 1, 2, 4, and 8 threads, and measure the network throughput performance between two nodes; network latency can be measured with the ping utility
- g. Compute the theoretical bandwidth and latency of your network. What efficiency do you achieve compared to the theoretical performance? Compare and contrast your performance to that achieved by iperf, and to the theoretical performance.
- h. Fill in the table 7 below for Network Throughput:

Proto-col	Con-currency	Block Size	MyNETBench Measured Throughput (Mb/sec)	iperf Measured Throughput (Mb/sec)	Theoretical Throughput (Mb/sec)	MyNETBench Efficiency (%)	iperf Efficiency (%)
TCP	1	1KB					
TCP	1	32KB					
TCP	2	1KB					
TCP	2	32KB					
TCP	4	1KB					
TCP	4	32KB					
TCP	8	1KB					
TCP	8	32KB					
UDP	1	1KB					
UDP	1	32KB					
UDP	2	1KB					
UDP	2	32KB					
UDP	4	1KB					
UDP	4	32KB					
UDP	8	1KB					
UDP	8	32KB					

- i. Fill in the table 8 below for Network Latency:

Proto-col	Con-currency	Message Size	MyNETBench Measured Latency (ms)	ping Measured Latency (ms)	Theoretical Latency (ms)	MyNETBench Efficiency (%)	iperf Efficiency (%)
TCP	1	1B					
TCP	2	1B					
TCP	4	1B					
TCP	8	1B					
UDP	1	1B					
UDP	2	1B					
UDP	4	1B					
UDP	8	1B					

Other requirements:

- You must write all benchmarks from scratch. You can use well known benchmarking software to verify your results, but you must implement your own benchmarks. Do not use code you find online, as you will get 0 credit for this assignment. If you have taken other courses where you

wrote similar benchmarks, you are welcome to start with your codebase as long as you wrote the code in your prior class.

- All of the benchmarks will have to evaluate concurrency performance; concurrency can be achieved using threads. Use strong scaling in all experiments, unless it is not possible, in which case you need to explain why a strong scaling experiment was not done. Be aware of the thread synchronizing issues to avoid inconsistency or deadlock in your system.
- Most benchmarks could be run on a single machine, but some benchmarks (e.g. network) will require 2 machines.
- Not all timing functions have the same accuracy; you must find one that has at least 1ms accuracy or better, assuming you are running the benchmarks for at least seconds at a time.
- Since there are many experiments to run, find ways (e.g. scripts) to automate the performance evaluation.
- For the best reliability in your results, repeat each experiment 3 times and report the average and standard deviation. This will help you get more stable results that are easier to understand and justify.
- No GUIs are required. Simple command line interfaces are expected.

2 Skeleton Code

- Skeleton code written in C for this assignment can be downloaded from a public git repo using this command.
`git clone ssh://<userid>@129.114.33.105/exports/git/public/cs553-pa1.git`
- If you prefer to use JAVA or Python, please make sure to follow the same naming conventions and input/output file formats as in the C code.
- Skeleton code has 4 folders – CPU, Disk, Memory and Network. Each folder has a C file with blank main function, a Makefile for compiling the code, input files and a sample output file format.
- ReadMe file in each folder has more details on the input format.
- Output should be space separated columns of values.
- Your main function should read input from the files provided and generate output in the format provided.

3 Where you will submit

You will have to submit your solution to a private git repository created for you. You will have to firstly clone the repository. Then you will have to add or update your source code, documentation and report. Your solution will be collected automatically after the deadline grace period. If you want to submit your homework later, you will have to push your final version of the solution and you will have let the TAs know of it through email cs553-s18@datasys.cs.iit.edu. There is no need to submit anything on BB for this assignment. Here are some examples on how to clone your private repository and how to add files to it (replace **userid** with your username):

```
git clone ssh://userid@129.114.33.105/exports/git/userid/cs553-
pa1.git
cd cs553-pa1/
touch readme.txt
cat "Username A20*" > readme.txt
```



```
git add readme.txt
git commit -m "Added the readme file"
git push
```

If you cannot access your repository contact the TAs. You can find a git cheat sheet here: <https://www.git-tower.com/blog/git-cheat-sheet/>

43 What you will submit

When you have finished implementing the complete assignment as described above, you should submit your solution to your private git repository. Each program must work correctly and be detailed in-line documented. You should hand in:

1. **Source code (30%):** All of the source code; in order to get full credit for the source code, your code must have in-line documents, must compile, and must be able to run the sample benchmarks from #2 above. You must have a makefile for easy compilation.
2. **Readme (10%):** A detailed manual describing how the program works. The manual should be able to instruct users other than the developer to run the program step by step. The manual should contain example commands to invoke each of the five benchmarks. This should be included as readme.txt in the source code folder.
3. **Report / Performance (60%):** Must have working code that compiles and runs on the specified cluster to receive credit for report/performance; furthermore, the code must match the performance presented in the report. A separate (typed) design document (named pa1-report.pdf) of approximately 1-3 pages describing the overall program design, and design tradeoffs considered and made. Also describe possible improvements and extensions to your program (and sketch how they might be made). Since this is an assignment aimed at teaching you about benchmarking, this is one of the most important part; you must evaluate the four benchmarks with the entire parameters space mentioned in Section 1, and put as a sub-section to your design document mentioned in (1) above. You must produce 8 tables to showcase the results; we encourage you to create additional figures to highlight your performance evaluation. Please combine data and plot on the same graph wherever possible, for either compactness reasons, or comparison reasons. Don't forget to plot the average and standard deviation (if you do multiple runs of an experiment). Also, you need to explain each graph's results in words. Hint: graphs with no axis labels, legends, well defined units, and lines that all look the same, are likely very hard to read and understand graphs. You will be penalized if your graphs are not clear to understand. Please specify which student contributed to what benchmark experiments.

Submit code through github.

Grades for late programs will be lowered 20% per day late beyond the 4-day late pass.